

# STA6106 Statistic Computing Project2

Robert Norberg, Jung-Han Wang

Wednesday, October 22, 2014

## Problem1

The function  $g$  is given by

$$g(x, y) = 4xy + (x + y^2)^2$$

The goal is to find the minimum of  $g$ .

### a. Minimize $g$ using Newton's method

The Newton's Algorithm is:

$$\mathbf{x}(n+1) = \mathbf{x}(n) - H(\mathbf{x}(n))^{-1} \nabla f(\mathbf{x}(n))$$

First step we need to verify that  $H(\mathbf{x}(n))$  is non-singular.

```
hx<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'x');hx ## Check H(X(n)) is a non-singular matrix
```

```
[1] 2
```

Setting up stopping condition to be:

$$\|\nabla f(\mathbf{x}(n))\| \leq \epsilon$$

```
newton <- function(f3, x0, tol = 1e-9, n.max = 100) {  
  # Newton's method starting at x0  
  # f3 is a function that given x returns the list  
  # f(x), f'(x), Hessian f''(x)  
  x <- x0 ## Set initial value  
  f3.x <- f3(x) ## Set Input Function  
  n <- 0 ## Set first turn n<-0  
  while ((max(abs(f3.x[[2]])) > tol) & (n < n.max)) { ##Set Convergence Criteria. If f'(x) greater than tol  
    x <- x - solve(f3.x[[3]], f3.x[[2]]) ##Calculate f'(x)/f''(x)  
    f3.x <- f3(x)  
    n <- n + 1 ##Continue to the next n  
  }  
  if (n == n.max) { ##If n=maximum value, output "newton failed to converge"  
    cat('newton failed to converge\n')  
  } else {  
    return(x)  
  }  
}
```

Then Calculate first and second derivatives for function  $g(x, y) = 4xy + (x + y^2)^2$  Using package (Ryacas).

```
library(Ryacas);
k <- function(x) {
4*x[1]*x[2]+(x[1]+x[2]^2)^2
}
##Calculate First Derivative of Function to X (f1)
f1<-D(expression(4*x*y+(x+y^2)^2), 'x');f1
```

4 \* y + 2 \* (x + y^2)

```
##Calculate Second Derivative of Function to X (f11)
f11<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'x');f11
```

[1] 2

```
##Calculate First Derivative of Function to y (f2)
f2<-D(expression(4*x*y+(x+y^2)^2), 'y');f2
```

4 \* x + 2 \* (2 \* y \* (x + y^2))

```
##Calculate Second Derivative of Function to X (f22)
f22<-D(D(expression(4*x*y+(x+y^2)^2), 'y'),'y');f22
```

2 \* (2 \* (x + y^2) + 2 \* y \* (2 \* y))

```
##Calculate Second Derivative of Function to X,y (f12)
f12<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'y');f12
```

4 + 2 \* (2 \* y)

Set up function with return value using  $f$ (original function),  $\nabla f$ (first derivative function) and  $H(x(n))$ (second derivative function).

```
f3 <- function(x) {
f <- 4*x[1]*x[2]+(x[1]+x[2]^2)^2
##Calculate First Derivative of Function to x[1] (f1)
f1<-4*x[2]+2*(x[1] + x[2]^2)
##Calculate Second Derivative of Function to x[1] (f11)
f11<-2
##Calculate First Derivative of Function to x[2] (f2)
f2<-4 * x[1] + 2 * (2 * x[2] * (x[1] + x[2]^2))
##Calculate Second Derivative of Function to x[1] (f22)
f22<-2 * (2 * (x[1] + x[2]^2) + 2 * x[2] * (2 * x[2]))
##Calculate Second Derivative of Function to x[1],x[2] (f12)
f12<-4 + 2 * (2 * x[2])
return(list(f, c(f1, f2), matrix(c(f11, f12, f12, f22), 2, 2))) ##Return 3 values, f, f'(x), f''(x)
}
```

Run the program with start point (0,0) and try add 0.5 each run. Left side of the result represents the value of starting point. Right side of the result represent the value of the ending point.

Find minimum point (0.8888889,-0.6666667), after try out the value of f3(X).

```

for (x0 in seq(0,1, .5)) {
  for (y0 in seq(0,1, .5)) {
    cat(c(x0,y0), '--(Left=Start Point, Right=Extreme Value)--', newton(f3, c(x0,y0)), '\n')
  }
}

```

```

0 0 --(Left=Start Point, Right=Extreme Value)-- 0 0
0 0.5 --(Left=Start Point, Right=Extreme Value)-- -7.439e-11 6.287e-11
0 1 --(Left=Start Point, Right=Extreme Value)-- 0 0
0.5 0 --(Left=Start Point, Right=Extreme Value)-- 0 0
0.5 0.5 --(Left=Start Point, Right=Extreme Value)-- -7.241e-16 6.847e-16
0.5 1 --(Left=Start Point, Right=Extreme Value)-- -4.226e-14 4.196e-14
1 0 --(Left=Start Point, Right=Extreme Value)-- 0 0
1 0.5 --(Left=Start Point, Right=Extreme Value)-- -1.311e-14 1.294e-14
1 1 --(Left=Start Point, Right=Extreme Value)-- 0.8889 -0.6667

```

```
f3(c(0.8888889,-0.6666667))
```

```

[[1]]
[1] -0.5926

```

```

[[2]]
[1] -2.222e-08 -2.815e-07

```

```

[[3]]
      [,1] [,2]
[1,] 2.000 1.333
[2,] 1.333 8.889

```

```
f3(c(-7.438748e-11,6.286889e-11))
```

```

[[1]]
[1] -1.317e-20

```

```

[[2]]
[1] 1.027e-10 -2.975e-10

```

```

[[3]]
      [,1] [,2]
[1,] 2 4.000e+00
[2,] 4 -2.975e-10

```

```
f3(c(-4.225566e-14,4.196096e-14 ))
```

```

[[1]]
[1] -5.307e-27

```

```

[[2]]
[1] 8.333e-14 -1.690e-13

```

```

[[3]]
      [,1] [,2]
[1,] 2 4.00e+00
[2,] 4 -1.69e-13

```

```
f3(c(-1.3112e-14,1.293993e-14))
```

```
[[1]]
```

```
[1] -5.067e-28
```

```
[[2]]
```

```
[1] 2.554e-14 -5.245e-14
```

```
[[3]]
```

```
      [,1]      [,2]  
[1,]      2 4.000e+00  
[2,]      4 -5.245e-14
```

**b. Minimize g using the steepest descent method. Use (1,0) as starting point.**

Define function using Golden-Section method to find

$$x_l < x_m < x_r$$

such that

$$f(x_l) \geq f(x_m) \text{ and } f(x_r) \geq f(x_m)$$

```
gsection = function(ftn, x.l, x.r, x.m, tol = 1e-9) {  
  # applies the golden-section algorithm to minimize ftn  
  # we assume that ftn is a function of a single variable  
  # and that x.l < x.m < x.r and ftn(x.l), ftn(x.r) >= ftn(x.m)  
  
  # the algorithm iteratively refines x.l, x.r, and x.m and terminates  
  # when x.r - x.l <= tol, then returns x.m  
  
  # golden ratio plus one  
  gr1 = 1 + (1 + sqrt(5))/2  
  # successively refine x.l, x.r, and x.m  
  f.l = ftn(x.l)  
  f.r = ftn(x.r)  
  f.m = ftn(x.m)  
  while ((x.r - x.l) > tol) {  
    if ((x.r - x.m) > (x.m - x.l)) {  
      y = x.m + (x.r - x.m)/gr1  
      f.y = ftn(y)  
      if (f.y <= f.m) {  
        x.l = x.m  
        f.l = f.m  
        x.m = y  
        f.m = f.y  
      } else {  
        x.r = y  
        f.r = f.y  
      }  
    } else {  
      y = x.m - (x.m - x.l)/gr1  
      f.y = ftn(y)  
      if (f.y <= f.m) {
```

```

        x.r = x.m
        f.r = f.m
        x.m = y
        f.m = f.y
    } else {
        x.l = y
        f.l = f.y
    }
}
}
return(x.m)
}

```

Create Function  $F=4xy + (x + y^2)^2$

```

##Set Function F
f <- function(x) {
f <- 4*x[1]*x[2]+(x[1]+x[2]^2)^2
  return(f)
}

```

Set up first order partial derivative Function F for x and y.

```

##Set Function Gradient F f'(x)
gradf<- function (x)
{##Calculate First Derivative of Function to x[1] (f1)
f1<-4*x[2]+2*(x[1] + x[2]^2)
##Calculate First Derivative of Function to x[2] (f2)
f2<-4 * x[1] + 2 * (2 * x[2] * (x[1] + x[2]^2))
  return(c(f1, f2))
}

```

Set up algorithm for gradient line search. Trying to find Minimum point between  $x_l$  and  $x_r$ .

```

line.search <- function(f, x, gradf, tol = 1e-9, a.max = 100) {
# x and gradf are vectors of length d
# g(a) = f(x + a*gradf) has a local minimum at a,
# within a tolerance
# if no local minimum is found then we use 0 or a.max for a
# the value returned is x + a*y
if (sum(abs(gradf)) == 0) return(x) # g(a) constant
g <- function(a) return(f(x + a*gradf))

# find a.l < a.m < a.r such that
# g(a.m) >= g(a.l) and g(a.m) >= g(a.r)
# a.l
a.l <- 0
g.l <- g(a.l)
# a.m
a.m <- 1
g.m <- g(a.m)
while ((g.m > g.l) & (a.m > tol)) {
a.m <- a.m/2

```

```

g.m <- g(a.m)
}
# if a suitable a.m was not found then use 0 for a
if ((a.m <= tol) & (g.m >= g.l)) return(x)
# a.r
a.r <- 2*a.m
g.r <- g(a.r)
while ((g.m >= g.r) & (a.r < a.max)) {
  a.m <- a.r
  g.m <- g.r
  a.r <- 2*a.m
  g.r <- g(a.r)
}
# if a suitable a.r was not found then use a.max for a
if ((a.r >= a.max) & (g.m > g.r)) return(x - a.max*gradf)
# apply golden-section algorithm to g to find a
a <- gsection(g, a.l, a.r, a.m)
return(x - a*gradf)
}

```

Set up iteration using Steepest Descent Algorithm.

```

descent <- function(f,gradf, x0, tol = 1e-9, n.max = 100) {
  # steepest descent algorithm
  # find a local minimum of f starting at x0
  # function gradf is the gradient of f
  x <- x0
  x.old <- x
  x <- line.search(f, x, gradf(x))
  n <- 1
  while (f(x.old)-(f(x)> tol) & (n < n.max)) {
    x.old <- x
    x <- line.search(f, x, gradf(x))
    n <- n + 1
  }
  return(x)
}

```

The minimize result is (0.8888889,-0.6666667), which is identical point to problem 1-1.

```
descent(f,gradf,c(1,0) )
```

```
[1] 0.8889 -0.6667
```

## Problem2

In 1986, the space shuttle Challenger exploded during takeoff, killing the seven astronauts aboard. The explosion was the result of an O-ring failure, a splitting of a ring of rubber that seals the parts of the ship together. The accident was believed to have been caused by the unusually cold weather ( $31^{\circ}F$  or  $0^{\circ}C$ ) at the time of launch, as there is reason to believe that the O-ring failure probabilities increase as temperature decreases. Data on previous space shuttle launches and O-ring failures is given in the dataset challenger

provided with the “mcsn” package of R. The first column corresponds to the failure indicators  $y_i$  and the second column to the corresponding temperature  $x_i$ , ( $1 \leq i \leq 24$ ).

```
# load challenger data
library(mcsn)
data(challenger)
```

a) The goal is to obtain MLEs for  $\beta_0$  and  $\beta_1$  in the following logistic regression model

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x$$

where  $p$  is the probability that at least one O-ring is damaged and  $x$  is the temperature. Create computer programs using Newton-Raphson algorithm to find MLEs of  $\hat{\beta}_0, \hat{\beta}_1$

The log-likelihood is given by

$$l(\beta) = \mathbf{y}^T \mathbf{Z} \beta - \mathbf{b}^T \mathbf{1}$$

where  $\mathbf{1}$  is a column vector of ones,  $\mathbf{y} = (y_1 \dots y_n)^T$ ,  $\mathbf{b} = (b(\theta_1) \dots b(\theta_n))^T$ , and  $\mathbf{Z}$  is the  $n \times 2$  matrix whose  $i$  th row is  $\mathbf{z}_i^T$ .

The score function is

$$\mathbf{l}'(\beta) = \mathbf{Z}^T (\mathbf{y} - \boldsymbol{\pi})$$

where  $\boldsymbol{\pi}$  is a column vector of the Bernoulli probabilities  $\pi_1, \dots, \pi_n$ . The Hessian is given by

$$\mathbf{l}''(\beta) = \frac{d}{d\beta} (\mathbf{Z}^T (\mathbf{y} - \boldsymbol{\pi})) = - \left( \frac{d\boldsymbol{\pi}}{d\beta} \right)^T \mathbf{Z} = -\mathbf{Z}^T \mathbf{W} \mathbf{Z}$$

where  $\mathbf{W}$  is a diagonal matrix with  $i$  th diagonal entry equal to  $\pi_i(1 - \pi_i)$ .

Newton's update is therefore

$$\beta^{(t+1)} = \beta^{(t)} - \mathbf{l}''(\beta^{(t)})^{-1} \mathbf{l}'(\beta^{(t)}) = \beta^{(t)} + \left( \mathbf{Z}^T \mathbf{W}^{(t)} \mathbf{Z} \right)^{-1} \left( \mathbf{Z}^T (\mathbf{y} - \boldsymbol{\pi}^{(t)}) \right)$$

where  $\boldsymbol{\pi}^{(t)}$  is the value of  $\boldsymbol{\pi}$  corresponding to  $\beta^{(t)}$ , and  $\mathbf{W}^{(t)}$  is the diagonal weight matrix evaluated at  $\boldsymbol{\pi}^{(t)}$ .

To find MLEs of  $\hat{\beta}_0$  and  $\hat{\beta}_1$  we designate the vector  $\mathbf{y}$  to be the response vector in the data and  $\mathbf{Z}$  to be the  $n \times 2$  matrix of ones and temperatures.

```
y <- challenger$oring
x <- challenger$temp
Z <- cbind(rep(1, nrow(challenger)), challenger$temp)
```

$\mathbf{Z}$  will remain unchanged, but for each iteration we will need to compute the predicted probabilities  $\pi_1, \dots, \pi_n$  using the current values of  $\hat{\beta}_0$  and  $\hat{\beta}_1$ . We will also need to compute for each iteration  $\mathbf{W}$ , the diagonal matrix with diagonal entry  $i$  equal to  $\pi_i(1 - \pi_i)$ . Below, we define a function to do each of these computations.

```

find_pis <- function(beta, x){
  # beta a vector of parameter estimates, x a vector of independent variable observations
  pis <- exp(beta%*%t(Z))/(1+exp(beta%*%t(Z)))
  return(pis)
}

find_W <- function(pis){
  W <- diag(c(pis*(1-pis)))
  return(W)
}

```

Next we define a function to calculate Newton's update using  $\mathbf{Z}$ ,  $\mathbf{y}$ ,  $\boldsymbol{\beta}$ , and the  $\pi$ s and  $\mathbf{W}$  calculated using the functions above.

```

find_update <- function(beta, Z, W, y, pis){
  numerator <- t(Z)%*%(y-t(pis))
  denominator <- t(Z)%*%W%*%Z
  update <- solve(denominator)%*%numerator
  return(beta+c(update))
}

```

Finally, we must define a stopping criterion so that the algorithm ceases to update at some point and returns to us values for  $\hat{\beta}_0$  and  $\hat{\beta}_1$ . We will stop updating our estimates when an iteration takes place that does not change them. We are not concerned with accuracy at more than seven decimal places, but we will write the program so that this may be changed.

We choose starting values for  $\hat{\beta}_0$  and  $\hat{\beta}_1$  to be  $[0, 0]$ , as is recommended in the textbook. Then we begin a while loop to update these estimates until convergence occurs.

```

beta <- c(0, 0)
old_beta <- c(NA, NA)
counter <- 0
decimal_places_of_precision <- 7
verbose <- TRUE

while(!identical(beta, old_beta)){
  pis <- find_pis(beta, x)
  W <- find_W(pis)
  # set old beta equal to current beta before updating
  old_beta <- beta
  # replace beta with updated estimates
  beta <- find_update(beta, Z, W, y, pis)
  # round to the number of decimal places desired
  beta <- round(beta, decimal_places_of_precision)
  counter <- counter+1
  # show us the current value of beta if "verbose" is TRUE
  if(verbose==T){
    msg <- paste0('Iteration ', counter, ': Beta = (', beta[1], ', ', beta[2], ')')
    cat(msg)
    cat('\n')
  }
}

```

Iteration 1: Beta = (9.6190476, -0.1495238)



```

Iteration 2: Beta = (13.6557372, -0.211247)
Iteration 3: Beta = (14.9382896, -0.2306001)
Iteration 4: Beta = (15.0422911, -0.2321537)
Iteration 5: Beta = (15.0429016, -0.2321627)
Iteration 6: Beta = (15.0429016, -0.2321627)

```

Now that the while loop has terminated, the current values of `beta` should be the MLEs of  $\hat{\beta}_0$  and  $\hat{\beta}_1$ . We check these against the estimates given using the `glm()` function in R.

```

# see our estimates
print(beta)

```

```

[1] 15.0429 -0.2322

```

```

# see the glm estimates
glm(y~x, family=binomial)$coefficients

```

```

(Intercept)          x
      15.0429      -0.2322

```

Our estimates match those produced by the `glm()` function, so we are confident that we have correctly found the MLEs of  $\hat{\beta}_0$  and  $\hat{\beta}_1$ .

**b) Solve the same problem using the “Iterative Reweighted Least Squares” algorithm and the Newton-Raphson algorithm to find MLEs of  $\hat{\beta}_0, \hat{\beta}_1$**

If we let

$$\mathbf{e}^{(t)} = \mathbf{y} - \boldsymbol{\pi}^{(t)}$$

and

$$\mathbf{x}^{(t)} = \mathbf{Z}\boldsymbol{\beta}^{(t)} + (\mathbf{W}^{(t)})^{-1}\mathbf{e}^{(t)}$$

then the Fisher scoring update can be written as

$$\begin{aligned}\boldsymbol{\beta}^{(t+1)} &= \boldsymbol{\beta} + \left(\mathbf{Z}^T \mathbf{W}^{(t)} \mathbf{Z}\right)^{-1} \mathbf{Z}^T \mathbf{e}^{(t)} \\ &= \left(\mathbf{Z}^T \mathbf{W}^{(t)} \mathbf{Z}\right)^{-1} \mathbf{Z}^T \mathbf{W}^{(t)} \mathbf{x}^{(t)}\end{aligned}$$

where  $\mathbf{x}^{(t)}$  is the *working response* that gets updated with each iteration.

If we reuse the functions defined earlier to calculate  $\boldsymbol{\pi}$  and  $\mathbf{W}$ , we just need to define a function for calculating  $\mathbf{x}$  and define a new function for updating using the iteratively reweighted least squares update before we begin to iterate. We define a function for calculating  $\mathbf{x}$  as follows:

```

find_x <- function(y, beta, pis, Z, W){
  e_t <- y - c(pis)
  x <- Z%*%beta + solve(W)%*%e_t
  return(x)
}

```

And a new updating function:

```
find_update_irls <- function(Z, W, working_response){
  numerator <- t(Z)%*%W%*%working_response
  denominator <- t(Z)%*%W%*%Z
  update <- solve(denominator)%*%numerator
  return(update)
}
```

Next we choose (0, 0) for the initial values of  $\beta$  and begin another while loop, this time updating using the iteratively reweighted least squares update. (We are still satisfied with 7 decimal places of accuracy.)

```
beta <- c(0, 0)
old_beta <- c(NA, NA)
counter <- 0
decimal_places_of_precision <- 7
verbose <- TRUE

while(!identical(beta, old_beta)){
  pis <- find_pis(beta, x)
  W <- find_W(pis)
  working_response <- find_x(y, beta, pis, Z, W)
  # set old beta equal to current beta before updating
  old_beta <- beta
  # replace beta with updated estimates
  beta <- find_update_irls(Z, W, working_response)
  # round to the number of decimal places desired
  beta <- round(c(beta), decimal_places_of_precision)
  counter <- counter+1
  # show us the current value of beta if "verbose" is TRUE
  if(verbose==T){
    msg <- paste0('Iteration ', counter, ': Beta = (', beta[1], ', ', beta[2], ')')
    cat(msg)
    cat('\n')
  }
}
```

```
Iteration 1: Beta = (9.6190476, -0.1495238)
Iteration 2: Beta = (13.6557372, -0.211247)
Iteration 3: Beta = (14.9382896, -0.2306001)
Iteration 4: Beta = (15.0422911, -0.2321537)
Iteration 5: Beta = (15.0429016, -0.2321627)
Iteration 6: Beta = (15.0429016, -0.2321627)
```

We arrive at the same MLEs as before, in the same number of iterations as before.

c) We are also interested in predicting O-ring failure. Challenger was launched at  $31^{\circ}F$ . What is the predicted probability of O-ring damage at  $31^{\circ}F$  ? How many O-ring failures should be expected at  $31^{\circ}F$  ? What can you conclude?

At  $31^{\circ}F$  we predict the probability of at least one O-ring failure to be

$$p = \frac{e^{\beta_0 + (\beta_1)(31)}}{1 + e^{\beta_0 + (\beta_1)(31)}} = \frac{e^{15.0429 + (-0.2322)(31)}}{1 + e^{15.0429 + (-0.2322)(31)}} = 0.9996$$

This prediction means that we are almost certain that at least one O-ring will fail at  $31^\circ F$ .

### Problem 3

The elastic net (Zou and Hastie, 2006) is considered to be a compromise between the ridge and lasso penalties. The elastic net can be formulated using the Lagrangian as follows:

$$\hat{\beta}^{enet} = \underset{\beta}{\operatorname{argmin}} \sum_{i=1}^n (y_i - x_i' \beta)^2 + \lambda_1 \sum_{j=1}^p |\beta_j| + \lambda_2 \sum_{j=1}^p \beta_j^2$$

where  $\lambda_1 \geq 0$  and  $\lambda_2 \geq 0$ . The “credit” data set is discussed in the textbook of James et al., p83. We will fit the elastic net model to the “credit” data set using only the quantitative predictors. Our challenge is to select the appropriate  $\lambda_1$  and  $\lambda_2$  before fitting the final model.

**a) Write a function in R using the cross-validation approach to find the optimum values of  $\lambda_1$  and  $\lambda_2$**

We opt to use ten fold cross-validation, but we will write the code so that it may be generalized for k-fold cross validation.

The first task for our ten fold cross-validation is to divide the data into ten folds. We borrow some code from last assignment to accomplish this.

```
k <- 10 # for 10-fold cross-validation
n <- nrow(credit) # sample size
seg_size <- floor(n/k) # approx size of each segment
leftovers <- n%%k # n mod k
seg_assignments <- rep(1:k, each=seg_size) # create segment assignments
if(leftovers > 0){
  seg_assignments <- c(seg_assignments, 1:leftovers) # add "leftover" observations if necessary
}
set.seed(8675309) # set random seed (is it bad that I always use 867-5309?)
seg_assignments <- sample(seg_assignments) # randomly permute segment assignments
```

For each value of  $k$ ,  $\lambda_1$ , and  $\lambda_2$  we want to assess the accuracy of the model somehow. We will use as our criterion the cross-validation error

$$CV_k(\lambda_1, \lambda_2) = \frac{1}{n_k} \sum_{i \in F_k} (y_i - \hat{f}_{(\lambda_1, \lambda_2)}^{-k}(x_i))^2$$

where  $n_k$  is the fold size,  $F_k$  is the set of observation in fold  $k$  (the test set), and  $\hat{f}_{(\lambda_1, \lambda_2)}^{-k}(x)$  is the prediction function from the model fit to all observations *not* in  $F_k$  using parameters  $(\lambda_1, \lambda_2)$ .

Next, we must conduct a grid search across many possible values of  $\lambda_1$  and  $\lambda_2$ . The `glmnet()` function in the `glmnet` package takes as arguments `alpha` and `lambda`, where what we are calling  $\lambda_1$  is given by  $\alpha\lambda$  and what we are calling  $\lambda_2$  is given by  $(1 - \alpha)\lambda$ , so we actually want to grid search over possible values of `lambda` and `alpha` being passed to the `glmnet()` function. It is important to note that using this parameterization  $\lambda$  is constrained to be larger than 0 and  $\alpha$  is constrained to be between 0 and 1.

A brief look at the documentation for the `glmnet()` function (see `?glmnet`) reveals that the argument `alpha` takes a scalar value, but the argument `lambda` may take a vector of many  $\lambda$  values to try. The documentation states

Do not supply a single value for `lambda` ... . Supply instead a decreasing sequence of `lambda` values. `glmnet` relies on its warm starts for speed, and it's often faster to fit a whole path than compute a single fit.

So for the sake of speed we will create a vector of (decreasing) `lambda` values to pass to the `glmnet()` function with each call and repeat this for several values of `alpha`, supplied in separate calls. We will do this for each of our  $k$  (10) folds, computing the validation errors,  $CV_k(\alpha, \lambda)$ , as we go.

Since we will pass many `lambda` values to `glmnet()` at a time, but only one `alpha` value at a time to the function, it will simplify things to define a function to extract the validation error,  $CV_k(\alpha, \lambda)$ , for all of the `lambdas` supplied with one call to the `glmnet()` function.

```
get_val_errs <- function(alpha, lambdas, train_data, train_response, test_data, test_response){  
  # fit the elastic net for all lambda values  
  fit <- glmnet(train_data, train_response, alpha=alpha, lambda=lambdas)  
  # use the models fit to predict test data  
  preds <- predict(fit, newx=test_data)  
  # calculate the mean squared error in the test data for each value of lambda  
  MSEs <- apply(preds, 2, function(x) mean((test_response-x)^2))  
  # return a data.frame with alpha, lambdas, and corresponding MSEs  
  return(data.frame(alpha=rep(alpha, length(MSEs)), lambda=lambdas, MSE=MSEs))  
}
```

Next we define the grid of `alpha` and `lambda` values over which we will search. We have refined this search grid by some trial and error.

```
lambda_vec <- seq(40, 0, by=-1) # lambdas to try for each alpha  
alphas <- seq(0, 1, by=0.01) # alphas to try
```

We will search the grid defined above once for each of our  $k$  (10) cross-validation folds, so we preallocate an empty list to hold the results for each of our  $k$  folds.

```
# preallocate an empty list with a slot for each of the k validation folds  
results <- vector(mode='list', length=k)
```

Finally we search over the defined grid of `alpha` and `lambda` values  $k$  times, once for each cross-validation fold, and compute the mean squared error of cross-validation at each pair of values.

```
for(i in 1:k){ # k validation folds  
  # define training data matrix  
  x <- model.matrix(Balance~., subset(credit, seg_assignments!=i))[, -1]  
  # and the training response vector  
  y <- credit$Balance[seg_assignments!=i]  
  
  # define the test data matrix  
  x_test <- model.matrix(Balance~., subset(credit, seg_assignments==i))[, -1]  
  # and the test response vector  
  y_test <- credit$Balance[seg_assignments==i]
```

```

# succinctly loop over alphas, finding MSE for each lambda as we go
result_list <- sapply(alphas, get_val_errs, lambdas=lambda_vec, train_data=x, train_response=y, test_data=x, test_response=y)
# convert the list into a data.frame using rbind()
result_df <- Reduce(rbind, result_list)
# assign a field for k
result_df$k <- i
# deposit results into preallocated list
results[[i]] <- result_df
}

```

The object `results` is now a list with  $k$  elements and each element is a `data.frame` with fields `k`, `alpha`, `lambda`, and `MSE`. We combine these homogenous list elements into one `data.frame` and find the average MSE for each set of `alpha` and `lambda` values (averaging the  $k$  MSE values calculated for each pair).

```

# average the MSEs from all k folds
results_all <- Reduce(rbind, results)
results_agg <- aggregate(MSE~alpha+lambda, data=results_all, mean)

```

The lowest MSEs appear to coincide with higher values of  $\alpha$  and values of  $\lambda$  in the 15-30 range.

```

bestMSE <- results_agg[which.min(results_agg$MSE),]
bestMSE

```


```

      alpha lambda    MSE
157  0.55      1 26666

```

The smallest MSE occurs at  $\alpha = 0.55$ ,  $\lambda = 1$ .

Since we are searching over two parameters we can visualize the cross-validation errors for each `alpha`, `lambda` pair with a heat map.



b) Repeat the same question as in (a) but using now the *one-standard-error* (1 SE) rule cross validation.

We define the standard error of cross validation as follows:

$$SE(\alpha, \lambda) = \frac{\sqrt{\text{var}(CV_1(\alpha, \lambda), \dots, CV_k(\alpha, \lambda))}}{\sqrt{k}}$$

For a particular  $(\alpha, \lambda)$  pair, this is the standard deviation of the  $k$  computed MSEs  $(CV_1(\alpha, \lambda), \dots, CV_k(\alpha, \lambda))$  divided by  $\sqrt{k}$ . Recall that each of the MSEs  $(CV_1(\alpha, \lambda), \dots, CV_k(\alpha, \lambda))$  are stored in the list object `results_all`. From these we can compute the standard error

$$SE(\alpha, \lambda)$$

for each  $(\alpha, \lambda)$  pair.

```
# calculate std dev of MSEs
results_SEs <- aggregate(MSE~alpha+lambda, data=results_all, sd)
# divide by sqrt(k)
results_SEs$MSE <- results_SEs$MSE/sqrt(k)
# rename column for less confusion
names(results_SEs) <- gsub('MSE', 'StdErr', names(results_SEs))
```

The standard error of the best  $(\alpha, \lambda)$  pair is

```
SE_best_model <- results_SEs[results_SEs$alpha==bestMSE$alpha & results_SEs$lambda==bestMSE$lambda, ]
SE_best_model
```

```
alpha lambda StdErr
157 0.55      1 3156
```


From this we compute the largest tolerable  $CV(\alpha, \lambda)$ .

```
MSE_max_tolerable <- bestMSE$MSE+SE_best_model$StdErr
```

Next, we retrieve all  $(\alpha, \lambda)$  pairs with a cross-validation error  $CV(\alpha, \lambda)$  less than this amount.

```
candidate_lambdas <- results_agg[results_agg$MSE<MSE_max_tolerable, c('alpha', 'lambda')]
```

These candidate  $(\alpha, \lambda)$  pairs are highlighted (in white) on the figure from earlier.



It turns out that there are many candidate models with a cross-validation error  $CV(\alpha, \lambda)$  below the threshold  $CV(\hat{\alpha}, \hat{\lambda}) + SE(\hat{\alpha}, \hat{\lambda})$ . In fact, we could choose literally any value of  $\alpha$  and any  $\lambda$  less than 20. We want to choose the simplest or most regularized model in this candidate model space. Since  $\lambda$  is the penalty parameter for model terms, larger values of  $\lambda$  will result in a simpler model. When  $\alpha = 1$ , we have the LASSO model, which results in simpler models by setting coefficient values ( $\beta$ 's) to 0. From the figure we can see that favoring  $\alpha = 1$  limits the largest candidate  $\lambda$  we can choose.

One way to define the “simplest model” is that which has the fewest nonzero parameters. Since some parameter values may be very small, but not quite zero, we will refine this definition slightly so that the “simplest model” is the model in the candidate model space with the fewest parameters larger in absolute value than 0.001. Given this definition, we fit models in the candidate model space on *all* of the data and select the model with the fewest parameter estimates  $|\hat{\beta}_p| \geq 0.001$ .

```
# define data matrix
x <- model.matrix(Balance~., credit)[-1]
# and the response vector
y <- credit$Balance

# define a function
find_nonzero_parameters <- function(alpha, lambda, data, response){
  fit <- glmnet(data, response, alpha=alpha, lambda=c(lambda))
  betas <- fit$beta@x # @ accesses a slot in an S4 object
  num_nonzero <- sum(abs(betas)>0.001)
  return(num_nonzero)
}
# use mapply() to iterate over alpha and lambda together
candidate_lambdas$num_params <- mapply(find_nonzero_parameters,
                                       alpha=candidate_lambdas$alpha,
                                       lambda=candidate_lambdas$lambda,
                                       MoreArgs=list(data=x, response=y))

# how many params does the smallest model have?
min(candidate_lambdas$num_params)
```

```
[1] 4
```

The model(s) with the fewest number of non-zero (approximately) parameters include four parameters. Several models in the candidate model space have four non-zero parameters however.

```
final_candidate_lambdas <- candidate_lambdas[candidate_lambdas$num_params==4,]
final_candidate_lambdas
```

	alpha	lambda	num_params
1616	1.00	15	4
1716	0.99	16	4
1717	1.00	16	4
1815	0.97	17	4
1816	0.98	17	4
1817	0.99	17	4
1818	1.00	17	4

Unsurprisingly, these all have large values of  $\alpha$ , which corresponds to the LASSO model, which we know tends to favor fewer parameters. Of these models, we will choose the one with the smallest cross-validation error  $CV(\alpha, \lambda)$ .

```
# merge with MSEs by alpha and lambda
final_candidate_lambdas <- merge(final_candidate_lambdas, results_agg, by=c('alpha', 'lambda'), all.x=T)
# best, simplest model using 1-SE rule
final_candidate_lambdas[which.min(final_candidate_lambdas$MSE),]
```

```
alpha lambda num_params MSE
5      1      15          4 29017
```

## Appendix with R code

```
# Clear working environment
rm(list=ls())
library(glmnet) # for fitting elastic net models
library(RCurl) # for getting data from the web
library(ggplot2) # for plots

# Options for document compilation
knitr::opts_chunk$set(warning=FALSE, message=FALSE, comment=NA, fig.width=4, fig.height=3)

#### Begin Problem 1 ####
hx<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'x');hx ## Check H(X(n)) is a non-singular matrix
newton <- function(f3, x0, tol = 1e-9, n.max = 100) {
  # Newton's method starting at x0
  # f3 is a function that given x returns the list
  # f(x), f'(x), Hessian f''(x)
  x <- x0 ## Set initial value
  f3.x <- f3(x) ## Set Input Function
  n <- 0 ## Set first turn n<-0
  while ((max(abs(f3.x[[2]])) > tol) & (n < n.max)) { ##Set Convergence Criteria. If f'(x) greater than t
  x <- x - solve(f3.x[[3]], f3.x[[2]]) ##Calculate f'(x)/f''(x)
  f3.x <- f3(x)
  n <- n + 1 ##Continue to the next n
  }
  if (n == n.max) { ##If n=maximum value, output "newton failed to converge"
  cat('newton failed to converge\n')
  } else {
  return(x)
  }
}

library(Ryacas);
k <- function(x) {
  4*x[1]*x[2]+(x[1]+x[2]^2)^2
}
##Calculate First Derivative of Function to X (f1)
f1<-D(expression(4*x*y+(x+y^2)^2), 'x');f1
##Calculate Second Derivative of Function to X (f11)
f11<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'x');f11
##Calculate First Derivative of Function to y (f2)
f2<-D(expression(4*x*y+(x+y^2)^2), 'y');f2
##Calculate Second Derivative of Function to X (f22)
```



```

f22<-D(D(expression(4*x*y+(x+y^2)^2), 'y'),'y');f22
##Calculate Second Derivative of Function to X,y (f12)
f12<-D(D(expression(4*x*y+(x+y^2)^2), 'x'),'y');f12
f3 <- function(x) {
  f <- 4*x[1]*x[2]+(x[1]+x[2]^2)^2
  ##Calculate First Derivative of Function to x[1] (f1)
  f1<-4*x[2]+2*(x[1] + x[2]^2)
  ##Calculate Second Derivative of Function to x[1] (f11)
  f11<-2
  ##Calculate First Derivative of Function to x[2] (f2)
  f2<-4 * x[1] + 2 * (2 * x[2] * (x[1] + x[2]^2))
  ##Calculate Second Derivative of Function to x[1] (f22)
  f22<-2 * (2 * (x[1] + x[2]^2) + 2 * x[2] * (2 * x[2]))
  ##Calculate Second Derivative of Function to x[1],x[2] (f12)
  f12<-4 + 2 * (2 * x[2])
  return(list(f, c(f1, f2), matrix(c(f11, f12, f12, f22), 2, 2))) ##Return 3 values, f, f'(x), f''(x)
}
for (x0 in seq(0,1, .5)) {
  for (y0 in seq(0,1, .5)) {
    cat(c(x0,y0), '--(Left=Start Point, Right=Extreme Value)--', newton(f3, c(x0,y0)), '\n')
  }
}

f3(c(0.8888889,-0.6666667))
f3(c(-7.438748e-11,6.286889e-11))
f3(c(-4.225566e-14,4.196096e-14 ))
f3(c(-1.3112e-14,1.293993e-14))

gsection = function(ftn, x.l, x.r, x.m, tol = 1e-9) {
  # applies the golden-section algorithm to minimize ftn
  # we assume that ftn is a function of a single variable
  # and that x.l < x.m < x.r and ftn(x.l), ftn(x.r) >= ftn(x.m)

  # the algorithm iteratively refines x.l, x.r, and x.m and terminates
  # when x.r - x.l <= tol, then returns x.m

  # golden ratio plus one
  gr1 = 1 + (1 + sqrt(5))/2
  # successively refine x.l, x.r, and x.m
  f.l = ftn(x.l)
  f.r = ftn(x.r)
  f.m = ftn(x.m)
  while ((x.r - x.l) > tol) {
    if ((x.r - x.m) > (x.m - x.l)) {
      y = x.m + (x.r - x.m)/gr1
      f.y = ftn(y)
      if (f.y <= f.m) {
        x.l = x.m
        f.l = f.m
        x.m = y
        f.m = f.y
      } else {
        x.r = y
        f.r = f.y
      }
    }
  }
  return(x.m)
}

```

```

    }
  } else {
    y = x.m - (x.m - x.l)/gr1
    f.y = ftn(y)
    if (f.y <= f.m) {
      x.r = x.m
      f.r = f.m
      x.m = y
      f.m = f.y
    } else {
      x.l = y
      f.l = f.y
    }
  }
}
}
return(x.m)
}

##Set Function F
f <- function(x) {
f <- 4*x[1]*x[2]+(x[1]+x[2]^2)^2
  return(f)
}

##Set Function Gradient F f'(x)
gradf<- function (x)
{##Calculate First Derivative of Function to x[1] (f1)
f1<-4*x[2]+2*(x[1] + x[2]^2)
##Calculate First Derivative of Function to x[2] (f2)
f2<-4 * x[1] + 2 * (2 * x[2] * (x[1] + x[2]^2))
  return(c(f1, f2))
}

line.search <- function(f, x, gradf, tol = 1e-9, a.max = 100) {
# x and gradf are vectors of length d
# g(a) =f(x +a*gradf) has a local minumum at a,
# within a tolerance
# if no local minimum is found then we use 0 or a.max for a
# the value returned is x + a*y
if (sum(abs(gradf)) == 0) return(x) # g(a) constant
g <- function(a) return(f(x + a*gradf))

# find a.l < a.m < a.r such that
# g(a.m) >=g(a.l) and g(a.m) >= g(a.r)
# a.l
a.l <- 0
g.l <- g(a.l)
# a.m
a.m <- 1
g.m <- g(a.m)
while ((g.m > g.l) & (a.m > tol)) {
a.m <- a.m/2
g.m <- g(a.m)
}

# if a suitable a.m was not found then use 0 for a
if ((a.m <= tol) & (g.m >= g.l)) return(x)

```

```

# a.r
a.r <- 2*a.m
g.r <- g(a.r)
while ((g.m >= g.r) & (a.r < a.max)) {
  a.m <- a.r
  g.m <- g.r
  a.r <- 2*a.m
  g.r <- g(a.r)
}
# if a suitable a.r was not found then use a.max for a
if ((a.r >= a.max) & (g.m > g.r)) return(x - a.max*gradf)
# apply golden-section algorithm to g to find a
a <- gsection(g, a.l, a.r, a.m)
return(x - a*gradf)
}
descent <- function(f,gradf, x0, tol = 1e-9, n.max = 100) {
  # steepest descent algorithm
  # find a local minimum of f starting at x0
  # function gradf is the gradient of f
  x <- x0
  x.old <- x
  x <- line.search(f, x, gradf(x))
  n <- 1
  while (f(x.old)-(f(x)> tol) & (n < n.max)) {
    x.old <- x
    x <- line.search(f, x, gradf(x))
    n <- n + 1
  }
  return(x)
}

descent(f,gradf,c(1,0) )

# #### Begin Problem 2 ####
# load challenger data
library(mcsm)
data(challenger)
y <- challenger$oring
x <- challenger$temp
Z <- cbind(rep(1, nrow(challenger)), challenger$temp)

find_pis <- function(beta, x){
  # beta a vector of parameter estimates, x a vector of independent variable observations
  pis <- exp(beta%*%t(Z))/(1+exp(beta%*%t(Z)))
  return(pis)
}

find_W <- function(pis){
  W <- diag(c(pis*(1-pis)))
  return(W)
}

find_update <- function(beta, Z, W, y, pis){

```

```

    numerator <- t(Z)%*%(y-t(pis))
    denominator <- t(Z)%*%W*%Z
    update <- solve(denominator)%*%numerator
    return(beta+c(update))
}
beta <- c(0, 0)
old_beta <- c(NA, NA)
counter <- 0
decimal_places_of_precision <- 7
verbose <- TRUE

while(!identical(beta, old_beta)){
  pis <- find_pis(beta, x)
  W <- find_W(pis)
  # set old beta equal to current beta before updating
  old_beta <- beta
  # replace beta with updated estimates
  beta <- find_update(beta, Z, W, y, pis)
  # round to the number of decimal places desired
  beta <- round(beta, decimal_places_of_precision)
  counter <- counter+1
  # show us the current value of beta if "verbose" is TRUE
  if(verbose==T){
    msg <- paste0('Iteration ', counter, ': Beta = (', beta[1], ', ', beta[2], ')')
    cat(msg)
    cat('\n')
  }
}
# see our estimates
print(beta)
# see the glm estimates
glm(y~x, family=binomial)$coefficients
find_x <- function(y, beta, pis, Z, W){
  e_t <- y-c(pis)
  x <- Z*%beta+solve(W)%*%e_t
  return(x)
}
find_update_irls <- function(Z, W, working_response){
  numerator <- t(Z)%*%W*%working_response
  denominator <- t(Z)%*%W*%Z
  update <- solve(denominator)%*%numerator
  return(update)
}
beta <- c(0, 0)
old_beta <- c(NA, NA)
counter <- 0
decimal_places_of_precision <- 7
verbose <- TRUE

while(!identical(beta, old_beta)){
  pis <- find_pis(beta, x)
  W <- find_W(pis)
  working_response <- find_x(y, beta, pis, Z, W)

```

```

# set old beta equal to current beta before updating
old_beta <- beta
# replace beta with updated estimates
beta <- find_update_irls(Z, W, working_response)
# round to the number of decimal places desired
beta <- round(c(beta), decimal_places_of_precision)
counter <- counter+1
# show us the current value of beta if "verbose" is TRUE
if(verbose==T){
  msg <- paste0('Iteration ', counter, ': Beta = (', beta[1], ', ', beta[2], ')')
  cat(msg)
  cat('\n')
}
}

##### Begin Problem 3 #####
myfile <- getURL("http://www-bcf.usc.edu/~gareth/ISL/Credit.csv") # grab the content of this page
credit <- read.csv(textConnection(myfile), header=T) # read the page content as a .csv
credit$X <- NULL # get rid of row ID column
# find only the column with a numeric class (or integer)
num_cols <- names(credit)[sapply(credit, class)%in%c('numeric', 'integer')]
# now subset the data set to only take the numeric columns
credit <- credit[, num_cols]
k <- 10 # for 10-fold cross-validation
n <- nrow(credit) # sample size
seg_size <- floor(n/k) # approx size of each segment
leftovers <- n%%k # n mod k
seg_assignments <- rep(1:k, each=seg_size) # create segment assignments
if(leftovers > 0){
  seg_assignments <- c(seg_assignments, 1:leftovers) # add "leftover" observations if necessary
}
set.seed(8675309) # set random seed (is it bad that I always use 867-5309?)
seg_assignments <- sample(seg_assignments) # randomly permute segment assignments
get_val_errs <- function(alpha, lambdas, train_data, train_response, test_data, test_response){
  # fit the elastic net for all lambda values
  fit <- glmnet(train_data, train_response, alpha=alpha, lambda=lambdas)
  # use the models fit to predict test data
  preds <- predict(fit, newx=test_data)
  # calculate the mean squared error in the test data for each value of lambda
  MSEs <- apply(preds, 2, function(x) mean((test_response-x)^2))
  # return a data.frame with alpha, lambdas, and corresponding MSEs
  return(data.frame(alpha=rep(alpha, length(MSEs)), lambda=lambdas, MSE=MSEs))
}
lambda_vec <- seq(40, 0, by=-1) # lambdas to try for each alpha
alphas <- seq(0, 1, by=0.01) # alphas to try
# preallocate an empty list with a slot for each of the k validation folds
results <- vector(mode='list', length=k)
for(i in 1:k){ # k validation folds
  # define training data matrix
  x <- model.matrix(Balance~., subset(credit, seg_assignments!=i))[, -1]
  # and the training response vector
  y <- credit$Balance[seg_assignments!=i]

```

```

# define the test data matrix
x_test <- model.matrix(Balance~., subset(credit, seg_assignments==i))[, -1]
# and the test response vector
y_test <- credit$Balance[seg_assignments==i]

# succinctly loop over alphas, finding MSE for each lambda as we go
result_list <- sapply(alphas, get_val_errs, lambdas=lambda_vec, train_data=x, train_response=y, test_
# convert the list into a data.frame using rbind()
result_df <- Reduce(rbind, result_list)
# assign a field for k
result_df$k <- i
# deposit results into preallocated list
results[[i]] <- result_df
}

# average the MSEs from all k folds
results_all <- Reduce(rbind, results)
results_agg <- aggregate(MSE~alpha+lambda, data=results_all, mean)
bestMSE <- results_agg[which.min(results_agg$MSE),]
bestMSE

# plot a heat map of the MSEs by alpha and lambda
ggplot(results_agg, aes(x=alpha, y=lambda, z=MSE))+
  geom_tile(aes(fill=MSE))+
  stat_contour(bins=12)+
  geom_point(data=bestMSE, color='red', size=3)+
  geom_text(data=bestMSE, aes(label=paste0('(', alpha, ', ', lambda, ')')), color='white', vjust=-1)+
  theme_bw(base_size=10)

# calculate std dev of MSEs
results_SEs <- aggregate(MSE~alpha+lambda, data=results_all, sd)
# divide by sqrt(k)
results_SEs$MSE <- results_SEs$MSE/sqrt(k)
# rename column for less confusion
names(results_SEs) <- gsub('MSE', 'StdErr', names(results_SEs))
SE_best_model <- results_SEs[results_SEs$alpha==bestMSE$alpha & results_SEs$lambda==bestMSE$lambda, ]

SE_best_model
MSE_max_tolerable <- bestMSE$MSE+SE_best_model$StdErr
candidate_lambdas <- results_agg[results_agg$MSE<MSE_max_tolerable, c('alpha', 'lambda')]
# plot a heat map of the MSEs by alpha and lambda
ggplot(results_agg, aes(x=alpha, y=lambda, z=MSE))+
  geom_tile(aes(fill=MSE))+
  stat_contour(bins=12)+
  geom_point(data=bestMSE, color='red', size=3)+
  geom_text(data=bestMSE, aes(label=paste0('(', alpha, ', ', lambda, ')')), color='white', vjust=-1)+
  geom_tile(data=candidate_lambdas, aes(fill=NULL, z=NULL), fill='white', alpha=0.5)+
  theme_bw(base_size=10)

# define data matrix
x <- model.matrix(Balance~., credit)[, -1]
# and the response vector
y <- credit$Balance

# define a function
find_nonzero_parameters <- function(alpha, lambda, data, response){
  fit <- glmnet(data, response, alpha=alpha, lambda=c(lambda))

```

```

betas <- fit$beta@x # @ accesses a slot in an s4 object
num_nonzero <- sum(abs(betas)>0.001)
return(num_nonzero)
}
# use mapply() to iterate over alpha and lambda together
candidate_lambdas$num_params <- mapply(find_nonzero_parameters,
                                       alpha=candidate_lambdas$alpha,
                                       lambda=candidate_lambdas$lambda,
                                       MoreArgs=list(data=x, response=y))
# how many params does the smallest model have?
min(candidate_lambdas$num_params)
final_candidate_lambdas <- candidate_lambdas[candidate_lambdas$num_params==4,]
final_candidate_lambdas
# merge with MSEs by alpha and lambda
final_candidate_lambdas <- merge(final_candidate_lambdas, results_agg, by=c('alpha', 'lambda'), all.x=T)
# best, simplest model using 1-SE rule
final_candidate_lambdas[which.min(final_candidate_lambdas$MSE),]

```

---